PRETEST

Section A: Background Information

1. **What is your current year of study?** (2nd year / 3rd year / 4th year / Postgraduate)
2. **Gender?** (Male / Female / Other / Prefer not to say)
3. **What is your current status?** (Full-time student / Working student)
4. **How would you rate your self-learning ability?** (Scale 1–5)
5. **Which is the first topic you choose for knowledge assessment?** (Dropdown 8 topics)
6. **Which is the second topic you choose for knowledge assessment?** (Dropdown 8 topics)
7. **Have you studied Topic 1 before?** (Never / Self-taught / University course / Online course)
8. **Have you studied Topic 2 before?** (Never / Self-taught / University course / Online course)
9. **What is your level of familiarity with Topic 1?** (Scale 1–5)
10. **What is your level of familiarity with Topic 2?** (Scale 1–5)
11. **Have you ever used GenAI tools (e.g., ChatGPT, Gemini) to support your learning?** (Yes / No)
12. **How often do you use AI for learning purposes?** (Daily / Several times a week / Rarely / Never)
13. **Which AI tool do you use most frequently?** (None / ChatGPT, Gemini, Copilot, Claude, Other...)
14. **What is your primary purpose for using AI in your studies?** (None / Solving exercises, explaining concepts, coding assistance, summarizing documents, etc.)
15. **Have you ever attended a course or workshop on how to use AI effectively?** (Yes / No)

Section B: Technical Knowledge Check

Association Rules Mining

Question	Bloom's Level	Question Content

1	Remembering	Which metric measures how frequently an itemset appears in the entire dataset? A. Confidence B. Support C. Lift
2	Remembering	The Apriori algorithm is primarily used for: A. Clustering similar data points B. Finding frequent itemsets C. Dimensionality reduction
3	Understanding	What does the Apriori Principle state about itemsets? A. If an itemset is frequent, all its subsets must also be frequent. B. If an itemset is frequent, its supersets must also be frequent. C. Only large item sets can be frequent.
4	Understanding	If a rule $\{A\} \rightarrow \{B\}$ has a Lift > 1, it implies that: A. A and B are independent. B. The presence of A increases the likelihood of B being present. C. A and B are negatively correlated.
5	Applying	In a dataset of 200 transactions, "Coke" appears in 100, "Chips" appears in 80, and both appear together in 40. Calculate the Confidence of $\{Coke\} \rightarrow \{Chips\}$. A. 0.2 B. 0.4 C. 0.5
6	Applying	Using the same data as Q5, calculate the Support for the itemset $\{Coke,\ Chips\}$. A. 20% B. 40% C. 50%

7	Analyzing	Why does the FP-Growth algorithm typically perform faster than Apriori on large datasets? A. It requires fewer database scans and no candidate generation. B. It uses a simpler mathematical formula for Lift. C. It only processes numerical data.
8	Analyzing	If you increase the minimum confidence threshold, how does it affect the discovered rules? A. It increases the number of rules. B. It decreases the number of rules but increases their reliability. C. It has no effect on the number of rules.
9	Evaluating	A retail manager wants to find products that are strongly associated but *not* just because they are both popular. Which metric should they prioritize? A. Support B. Confidence C. Lift
10	Creating	Suppose you are designing a market basket analysis for a pharmacy. How would you modify the traditional ARM process to prioritize "Critical Drug Interactions" over "Frequent Purchases"?

Recommender System

Question	Bloom's Level	Question Content

1	**Remembering**	Which type of recommender system suggests items similar to those a user liked in the past based on item features? A. Collaborative Filtering B. Content-based Filtering C. Association Rules
2	**Remembering**	What is the term for the situation where a system cannot recommend items to a new user due to a lack of history? A. Data Sparsity B. Overfitting C. Cold Start Problem
3	**Understanding**	How does User-User Collaborative Filtering differ from Item-Item Collaborative Filtering? A. User-User finds similar users; Item-Item finds similar items based on user ratings. B. User-User only uses metadata; Item-Item only uses ratings. C. There is no fundamental difference.
4	**Understanding**	Why is Matrix Factorization used in Recommender Systems? A. To increase the size of the dataset. B. To uncover latent features that explain the ratings in a sparse matrix. C. To categorize users into fixed demographic groups.
5	**Applying**	You are building a movie recommendation. If a user likes "Inception" (Sci-Fi, Nolan) and "Interstellar" (Sci-Fi, Nolan), which movie is a Content-based system most likely to recommend? A. A popular Romantic Comedy B. "Tenet" (Sci-Fi, Nolan) C. A random documentary

6	**Applying**	In a rating matrix, User A and User B have a high Cosine Similarity. If User A highly rates Movie X, what will the system likely predict for User B regarding Movie X? A. A low rating B. A high rating C. No rating possible
7	**Analyzing**	Consider a system with millions of items but only a few ratings per user. What is the most significant challenge this system faces? A. Computational speed B. Data Sparsity C. Hardware storage
8	**Analyzing**	Compare Precision and Recall in the context of recommendations. Which one is more important if the goal is to ensure the "Top-5" list contains only relevant items? A. Precision B. Recall C. Both are equally important
9	**Evaluating**	Evaluate the ethical implications of a "Filter Bubble" in news recommendation systems. How does it affect the diversity of information for the user?
10	**Creating**	Propose a Hybrid Recommender System architecture that effectively combines Collaborative Filtering and Content-based filtering to mitigate the Cold Start problem for new items.

Fuzzy Logic

Question	Bloom's Level	Question Content

1	Remembering	What is the essential range of values used for the degree of membership in a fuzzy set? A. Only 0 or 1 B. Any real number C. A value in the interval [0, 1]
2	Remembering	Which component of a Fuzzy Inference System (FIS) translates the fuzzy output into a single, definite crisp value? A. Fuzzifier B. Inference Engine C. Defuzzifier
3	Understanding	In the context of fuzzy set operations, how is the membership value for "Set A OR Set B" ($A \cup B$) typically calculated using the standard fuzzy Max operator? A. $\mu_{A \cup B}(x) = \text{Min}(\mu_A(x), \mu_B(x))$ B. $\mu_{A \cup B}(x) = \text{Max}(\mu_A(x), \mu_B(x))$ C. $\mu_{A \cup B}(x) = 1 - \mu_A(x)$
4	Understanding	Explain the primary limitation of using traditional Boolean (crisp) logic when dealing with human linguistic concepts like "Hot" or "Fast." A. It requires too many complex mathematical functions. B. It forces a binary distinction (either 0 or 1), failing to capture partial truth. C. It only works with discrete variables, not continuous ones.
5	Applying	A process uses the following rule: IF Pressure IS High (μ=0.9) AND Temperature IS Low (μ=0.5), THEN Valve IS Closed. Using the Min operator for AND, what is the 'firing strength' of this rule? A. 0.9 B. 0.5 C. 0.7
6	Applying	A fuzzy set 'Medium Speed' has a membership function defined by a trapezoid. If a speed of 50 km/h has a membership value of 1.0, and a speed of 70 km/h has a membership value of 0.5, what method is being used to assign these values? A. Expert knowledge (Heuristics)

		B. Random Assignment C. Statistical Regression
7	Analyzing	When designing a fuzzy controller, why is the selection and shape of the Membership Functions (e.g., triangular, trapezoidal, Gaussian) considered critical? A. They directly determine how the system interprets the input data and define the meaning of linguistic terms. B. They only affect the final defuzzification step. C. They are chosen randomly as long as they cover the input range.
8	Analyzing	Compare the Mamdani and Sugeno methods of Fuzzy Inference. What is the fundamental difference between their output definitions? A. Mamdani uses fuzzy sets as outputs; Sugeno uses crisp, mathematical functions as outputs. B. Mamdani is used only for classification; Sugeno is used for control. C. Mamdani is faster; Sugeno is more accurate.
9	Evaluating	A factory needs a system to control the chemical mix for a product. One system uses fixed, sharp thresholds. The other uses a Fuzzy Logic system that allows for smooth transitions between control states. Evaluate which system is likely to produce smoother and less abrupt control actions, and why.
10	Creating	Propose a simple linguistic rule for a Fuzzy Logic system designed to control a room thermostat, using the inputs "Room Temperature" (e.g., Cold, Perfect, Hot) and "Temperature Change Rate" (e.g., Falling, Stable, Rising).

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembering	C. A value in the interval [0, 1]
2	Remembering	C. Defuzzifier

3	Understand…	B. $\mu_{A \cup B}(x) = Max(\mu_A(x), \mu_B(x))$
4	Understand…	B. It forces a binary distinction (either 0 or 1), failing to capture partial truth.
5	Applying	B. 0.5 (Calculated as Min(0.9, 0.5))
6	Applying	A. Expert knowledge (Heuristics)
7	Analyzing	A. They directly determine how the system interprets the input data and define the meaning of linguistic terms.
8	Analyzing	A. Mamdani uses fuzzy sets as outputs; Sugeno uses crisp, mathematical functions as outputs.
9	Evaluating	The Fuzzy Logic system is likely to produce smoother and less abrupt control actions. This is because it uses continuous Membership Functions, allowing for gradual transitions in the control output as inputs change. The traditional system, with its fixed, sharp thresholds, would cause sudden, discontinuous changes when the input crosses a boundary.
10	Creating	Example Rule: IF Room Temperature IS Cold AND Temperature Change Rate IS Falling, THEN Heater Output IS Maximum (High).

Linear Regression

Question	Bloom's Level	Question Content
1	Remembering	What is the primary goal of Simple Linear Regression? A. To find the best fitting curve. B. To model the relationship between one independent variable (X) and one dependent variable (Y) using a straight line. C. To classify data points into discrete categories.
2	Remembering	In the linear regression equation $Y = \beta_0 + \beta_1 X + \epsilon$, what does β_1 represent? A. The Y-intercept.

		B. The slope of the regression line. C. The residual error.
3	Understanding	Explain the meaning of a coefficient of determination (R^2) value of 0.85. A. 85% of the data points fall directly on the regression line. B. 85% of the variance in the dependent variable (Y) is predictable from the independent variable (X). C. The model has an 85% chance of predicting the correct value.
4	Understanding	What is the fundamental difference between Simple Linear Regression and Multiple Linear Regression? A. Simple Linear Regression has continuous variables; Multiple Linear Regression has categorical variables. B. Simple Linear Regression uses one predictor variable; Multiple Linear Regression uses two or more predictor variables. C. Multiple Linear Regression is only used for non-linear relationships.
5	Applying	A model predicts house price (Y) based on square footage (X) with the equation Price = 50,000 + 100 · SqFt. If a house is 2,000SqFt, what is the predicted price? A. $250,000 B. $200,000 C. $150,000
6	Applying	Which loss function is most commonly minimized during the training of a standard Linear Regression model? A. Mean Absolute Error (MAE) B. Binary Cross-Entropy C. Mean Squared Error (MSE)
7	Analyzing	When performing residual analysis, if you observe a distinct pattern (e.g., a parabolic shape) in the plot of residuals versus fitted values, what does this indicate about the linear regression model? A. The model is highly accurate. B. The assumption of linearity has been violated. C. The data is perfectly correlated.

8	Analyzing	Why is it crucial to check for multicollinearity in Multiple Linear Regression, and what statistical measure is often used to detect it? A. Multicollinearity makes the model faster; VIF is used. B. Multicollinearity inflates the standard errors of the coefficients, making them unstable; Variance Inflation Factor (VIF) is used. C. Multicollinearity is only a problem in simple regression; R-squared is used.
9	Evaluating	Evaluate the trade-off between using a high-degree polynomial regression model (which perfectly fits the training data) versus a simple linear regression model for future prediction. Which model is generally preferred for generalization, and why?
10	Creating	Propose a hypothetical feature transformation (e.g., squaring, logging) for the independent variable X in a dataset where the relationship between X and Y is initially exponential, and briefly explain why this transformation is necessary for Linear Regression.

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembering	B. To model the relationship between one independent variable (X) and one dependent variable (Y) using a straight line.
2	Remembering	B. The slope of the regression line.
3	Understand…	B. 85% of the variance in the dependent variable (Y) is predictable from the independent variable (X).
4	Understand…	B. Simple Linear Regression uses one predictor variable; Multiple Linear Regression uses two or more predictor variables.
5	Applying	A. 250,000 (Calculated as $50,000 + 100 \cdot 2000 = 250,000$

6	Applying	C. Mean Squared Error (MSE)
7	Analyzing	B. The assumption of linearity has been violated.
8	Analyzing	B. Multicollinearity inflates the standard errors of the coefficients, making them unstable; Variance Inflation Factor (VIF) is used.
9	Evaluating	The simple linear regression model is generally preferred for generalization. The high-degree polynomial model, while perfectly fitting the training data, is prone to overfitting, which leads to poor performance on new, unseen data.
10	Creating	Example Transformation: Transform x to logX (logarithmic transformation). This is necessary because Linear Regression assumes a linear relationship between X and Y. The log transformation effectively linearizes the initial exponential relationship, allowing the standard Linear Regression assumptions to be met.

Logistic Regression

Question	Bloom's Level	Question Content
1	Remembering	What is the primary difference between Linear Regression and Logistic Regression regarding the dependent variable (Y)? A. Linear Regression predicts continuous values; Logistic Regression predicts probabilities for binary outcomes. B. Linear Regression uses categorical inputs; Logistic Regression uses continuous inputs. C. They both predict continuous values but use different optimization methods.
2	Remembering	Which function is applied in Logistic Regression to map the output of the linear equation into a probability range between 0 and 1?

		A. Rectified Linear Unit (ReLU) B. Softmax Function C. Sigmoid Function
3	Understanding	Why does Logistic Regression use the Sigmoid function instead of directly using the linear equation $Y = \beta_0 + \beta_1 X + \epsilon$? A. To ensure the predicted outcome is non-negative. B. To transform the output into a value that can be interpreted as a valid probability (between 0 and 1). C. To minimize the Mean Squared Error (MSE).
4	Understanding	In Logistic Regression, the log-odds (logit) is defined as the logarithm of the ratio of the probability of success (p) to the probability of failure (1 - p). What is the formula for the log-odds? A. $\ln(\frac{p}{1-p})$ B. p + (1-p) C. 1 / p
5	Applying	A Logistic Regression model provides a log-odds (logit) value of 0 for a specific customer. What is the calculated probability (p) that this customer belongs to the positive class? A. 0 B. 0.5 C. 1
6	Applying	Which mathematical function is typically minimized as the cost function (loss function) when training a Logistic Regression model? A. Mean Squared Error (MSE) B. R-squared C. Cross-Entropy Loss (or Log Loss)
7	Analyzing	How should a coefficient β be interpreted in Logistic Regression? A. A one-unit increase in X increases the predicted probability p by β.

		B. A one-unit increase in X increases the predicted Y by β. C. A one-unit increase in X increases the log-odds of the positive class by β.
8	Analyzing	In a classification dataset, if the data points are perfectly separated by a linear boundary, why might training a standard (unregularized) Logistic Regression model lead to numerical issues? A. The model will converge immediately. B. The coefficients (β) will approach infinity, indicating complete separation (perfect fit). C. The Cross-Entropy Loss will become zero, halting the training process prematurely.
9	Evaluating	For a system designed to predict whether a manufacturing machine will fail (positive class) or run normally (negative class), which metric should be prioritized: Precision (minimizing False Positives) or Recall (minimizing False Negatives), and why?
10	Creating	Propose a specific architecture or method to adapt a Binary Logistic Regression model for a problem with three or more mutually exclusive classes (e.g., predicting Low, Medium, or High risk levels).

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembe…	A. Linear Regression predicts continuous values; Logistic Regression predicts probabilities for binary outcomes.
2	Remembe…	C. Sigmoid Function
3	Understan…	B. To transform the output into a value that can be interpreted as a valid probability (between 0 and 1).
4	Understan…	A. $\ln(\frac{p}{1-p})$
5	Applying	B. 0.5 (If $\ln(\frac{p}{1-p}) = 0$, then $\frac{p}{1-p} = e^0 = 1$.

		Therefore, p = 1 - p, meaning 2p = 1, or p = 0.5).
6	Applying	C. Cross-Entropy Loss (or Log Loss)
7	Analyzing	C. A one-unit increase in X increases the log-odds of the positive class by β.
8	Analyzing	B. The coefficients (β) will approach infinity, indicating complete separation (perfect fit).
9	Evaluating	Recall (minimizing False Negatives) should be prioritized. A False Negative (missing a machine failure) is usually more costly and dangerous than a False Positive (predicting failure when the machine is fine), as failure can lead to severe operational damage or accidents.
10	Creating	The problem requires Multinomial Logistic Regression (also known as Softmax Regression), which generalizes the sigmoid function to handle multiple classes simultaneously, or using the One-vs-Rest (OvR) approach, which trains separate binary Logistic Regression models for each class against all others.

Latent Dirichlet Allocation

Question	Bloom's Level	Question Content
1	Remembering	**What is the primary purpose of Latent Dirichlet Allocation (LDA)?** A. Classifying documents into predefined categories B. Discovering hidden topics in a collection of documents C. Translating text between languages
2	Remembering	**In LDA, what is a “topic”?** A. A distribution over words B. A single keyword representing a document C. A cluster of documents

3	Understanding	**How does LDA represent a document?** A. As a single topic B. As a mixture of multiple topics C. As a list of keywords only
4	Understanding	**Why is LDA considered an "unsupervised learning" method?** A. It requires labeled training data B. It only works with structured datasets C. It learns patterns without labeled data
5	Applying	**You apply LDA to a collection of news articles. What is the expected output?** A. A classification label for each article B. A set of topics and their word distributions C. A sentiment score for each article
6	Applying	**If a document has 70% Topic A and 30% Topic B, what does this imply?** A. The document belongs only to Topic A B. The document contains words mostly related to Topic A but also some from Topic B C. The document is incorrectly modeled
7	Analyzing	**What happens if you set the number of topics (K) too high in LDA?** A. Topics become fragmented and less meaningful B. Topics become too general C. The model stops working
8	Analyzing	**Why is preprocessing (e.g., removing stopwords) important before applying LDA?** A. To reduce computation time only B. To improve topic quality by removing non-informative words C. It is not necessary for LDA

9	Evaluating	**You run LDA on a dataset and obtain topics with overlapping and unclear keywords. What does this indicate?** (Short answer) Expected idea: ● Poor parameter choice (K, alpha, beta) ● Poor preprocessing ● Dataset not suitable
10	Creating	**Propose one real-world application of LDA and explain how it would be useful.** (Open-ended) Expected examples: ● News clustering ● Recommendation systems ● Customer review analysis

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembe…	B. Discovering hidden topics in a collection of documents
2	Remembe…	A. A distribution over words
3	Understan…	B. As a mixture of multiple topics
4	Understan…	C. It learns patterns without labeled data
5	Applying	B. A set of topics and their word distributions
6	Applying	B. The document contains words mostly related to Topic A but also some from Topic B
7	Analyzing	A. Topics become fragmented and less meaningful

8	Analyzing	B. To improve topic quality by removing non-informative words
9	Evaluating	Mentions at least one valid reason: wrong K, poor preprocessing, bad parameters, or unsuitable data.
10	Creating	Provides one real-world application + brief correct explanation of how LDA is used.

Deep Neural Networks

Question	Bloom's Level	Question Content
1	Remembering	**What is a Deep Neural Network (DNN)?** A. A network with only one layer B. A neural network with multiple hidden layers C. A database system
2	Remembering	**What is the role of an activation function in a neural network?** A. Store data B. Introduce non-linearity C. Reduce dataset size
3	Understanding	**Why are multiple hidden layers useful in deep learning?** A. They reduce computation B. They allow the model to learn complex patterns C. They remove the need for data
4	Understanding	**What is overfitting in deep neural networks?** A. Model performs well on training data but poorly on new data B. Model cannot learn patterns C. Model is too simple

5	Applying	**You train a DNN and notice high training accuracy but low validation accuracy. What should you do?** A. Increase overfitting B. Apply regularization or dropout C. Remove training data
6	Applying	**If a neural network cannot learn complex patterns, what change might help?** A. Reduce layers B. Add more hidden layers or neurons C. Remove activation functions
7	Analyzing	**What happens if a network is too deep without proper techniques (e.g., normalization)?** A. Vanishing/exploding gradients may occur B. Model becomes faster C. Accuracy always increases
8	Analyzing	**Why is data normalization important for training DNNs?** A. It increases randomness B. It stabilizes and speeds up training C. It removes labels
9	Evaluating	**A DNN model performs poorly on both training and validation sets. What does this indicate?** (Short answer)
10	Creating	**Propose one real-world application of Deep Neural Networks and explain its usefulness.** (Open-ended)

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembe…	B. A neural network with multiple hidden layers
2	Remembe…	B. Introduce non-linearity
3	Understan…	B. They allow the model to learn complex patterns

4	Understan…	A. Model performs well on training data but poorly on new data
5	Applying	B. Apply regularization or dropout
6	Applying	B. Add more hidden layers or neurons
7	Analyzing	A. Vanishing/exploding gradients may occur
8	Analyzing	B. It stabilizes and speeds up training
9	Evaluating	Mentions underfitting / model too simple / insufficient capacity / poor training
10	Creating	Provides valid application + correct explanation of DNN usage

Word Embedding

Question	Bloom's Level	Question Content
1	Remembering	**What is a word embedding?** A. A method to store images B. A dense vector representation of words C. A type of database
2	Remembering	**Which model is commonly used for learning word embeddings?** A. Decision Tree B. Word2Vec C. K-Means
3	Understanding	**What does it mean if two words have similar embeddings?** A. They are identical words B. They have the same length C. They appear in similar contexts

4	Understanding	**What is the main difference between one-hot encoding and word embeddings?** A. One-hot is dense, embeddings are sparse B. One-hot captures semantic similarity, embeddings do not C. Embeddings capture semantic relationships, one-hot does not
5	Applying	**You want to find words with similar meanings in a dataset. What should you use?** A. One-hot encoding B. Word embeddings with cosine similarity C. Random vectors
6	Applying	**In Word2Vec, what is the goal of the Skip-gram model?** A. Predict the target word from context B. Predict context words from a target word C. Remove rare words
7	Analyzing	**Why do word embeddings often perform better than one-hot encoding in NLP tasks?** A. They capture semantic and syntactic relationships B. They reduce vocabulary size C. They eliminate the need for training
8	Analyzing	**What problem arises when the same word has multiple meanings (e.g., "bank")?** A. Overfitting B. Polysemy problem C. Vanishing gradient
9	Evaluating	**Word embeddings produce poor similarity results. What could be the reasons?** (Short answer)
10	Creating	**Propose one application of word embeddings and explain how they improve performance.** (Open-ended)

Answer Key

Question	Bloom's Level	Correct Answer / Explanation
1	Remembe…	B. A dense vector representation of words
2	Remembe…	B. Word2Vec
3	Understan…	C. They appear in similar contexts
4	Understan…	C. Embeddings capture semantic relationships, one-hot does not
5	Applying	B. Word embeddings with cosine similarity
6	Applying	B. Predict context words from a target word
7	Analyzing	A. They capture semantic and syntactic relationships
8	Analyzing	B. Polysemy problem
9	Evaluating	Mentions at least one: poor training data, small corpus, bad hyperparameters, lack of preprocessing
10	Creating	Provides valid application + correct explanation (e.g., search, recommendation, NLP tasks)

Section C: AI Attitudes & Readiness

Instructions: Please rate your level of agreement with the following statements on a scale from 1 to 5. *(1: Strongly Disagree | 2: Disagree | 3: Neutral | 4: Agree | 5: Strongly Agree)*

Part 1: Perceived Usefulness

1. I believe that AI can help me understand difficult concepts faster.
2. I think that using AI helps me save a significant amount of study time.
3. I find AI to be an effective assistant in summarizing lengthy or complex technical documents.
4. I believe that mastering AI tools is a mandatory skill for my future career in technology.

Part 2: Trust

5. I am generally confident in the accuracy of the technical explanations provided by AI.
6. I feel comfortable relying on the algorithmic logic and source code generated by AI.
7. I trust AI as a reliable resource when exploring new and unfamiliar technical topics.

Part 3: AI Learning Self-efficacy

8. I feel confident in my ability to use AI as a supportive tool for my self-study process.
9. I know how to craft effective prompts to obtain high-quality responses from AI.
10. I am capable of refining my queries if the initial response from the AI is not satisfactory.
11. I feel comfortable using AI to troubleshoot or debug technical problems in my assignments.

Part 4: Critical Awareness & Readiness

12. I can accurately identify when an AI-generated response is unreliable or inconsistent.
13. I believe it is essential for learners to verify AI-generated answers against official textbooks or sources.
14. I am aware that AI may occasionally provide incorrect information in a very convincing manner.
15. I am ready to use AI as a learning support tool specifically for the Intelligent Systems (IS) course.